Experiment 4: Ensemble Prediction and Tree-Based Model Evaluation with Hyperparameter Optimization

Machine Learning Lab Report

Academic Year 2025–2026

Aim

To implement Decision Tree, Random Forest, AdaBoost, Gradient Boosting, XGBoost, and Stacking Classifiers on the Wisconsin Breast Cancer dataset, optimize hyperparameters using GridSearchCV, and evaluate their performance with ROC curves, Confusion Matrices, and 5-Fold Cross Validation.

Libraries Used

- pandas, numpy, matplotlib, seaborn
- scikit-learn
- xgboost

1. Imports and Setup

```
# ================= IMPORTS =================
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import time
import warnings
warnings.filterwarnings("ignore")
```

```
9
10 from sklearn.model_selection import train_test_split, cross_val_score,
       KFold, GridSearchCV
11 from sklearn.preprocessing import StandardScaler, LabelEncoder
12 from sklearn.metrics import (
13     accuracy_score, precision_score, recall_score, f1_score,
14     confusion_matrix, roc_curve, auc, ConfusionMatrixDisplay
15 )
16
17 # Models
18 from sklearn.tree import DecisionTreeClassifier
19 from sklearn.ensemble import AdaBoostClassifier,
       GradientBoostingClassifier, RandomForestClassifier, StackingClassifier
20 from sklearn.naive_bayes import GaussianNB
21 from sklearn.svm import SVC
22 from sklearn.linear_model import LogisticRegression
23 from sklearn.neighbors import KNeighborsClassifier
24 import xgboost as xgb
```

2. Data Loading and Preprocessing

```
1 # ================= LOAD & PREPROCESS =================
2 cols = ["ID", "Diagnosis"] + [f"feature_{i}" for i in range(1, 31)]
3 df = pd.read_csv("wdbc.data", header=None, names=cols)
4
5 # Drop ID column
6 df.drop(columns=["ID"], inplace=True)
7
8 # Encode labels (M = malignant     1, B = benign     0)
9 df["Diagnosis"] = LabelEncoder().fit_transform(df["Diagnosis"])
10
11 # Features / Target
12 X_raw = df.drop(columns=["Diagnosis"])
13 y = df["Diagnosis"]
14
15 # Standardize
16 scaler = StandardScaler()
17 X_scaled = scaler.fit_transform(X_raw)
18
19 # Train/test split
20 X_train, X_test, y_train, y_test = train_test_split(
21     X_scaled, y, test_size=0.2, random_state=42, stratify=y
22 )
```

3. Exploratory Data Analysis (EDA)

```
1 # ================= EDA =================
```

```
sns.countplot(x=y)
plt.title("Class Balance (Benign=0, Malignant=1)")
plt.show()

plt.figure(figsize=(10, 8))
sns.heatmap(df.drop(columns=["Diagnosis"]).corr(), cmap="coolwarm")
plt.title("Feature Correlation Heatmap")
plt.show()
```

4. Evaluation Function (ROC + Confusion Matrix)

```
# ================ EVALUATION FUNCTION ================
def evaluate(name, model, X_test, y_test):
    fig, (ax1, ax2) = plt.subplots(1, 2, figsize=(12, 5))
    fig.suptitle(f'{name} Performance', fontsize=16)

    if hasattr(model, "predict_proba"):
        probs = model.predict_proba(X_test)[:, 1]
        fpr, tpr, _ = roc_curve(y_test, probs)
        roc_auc = auc(fpr, tpr)

        ax1.plot(fpr, tpr, color='darkorange', lw=2, label=f'ROC curve (
    AUC = {roc_auc:.2f})')
        ax1.plot([0, 1], [0, 1], color='navy', lw=2, linestyle='--')
        ax1.set_xlim([0.0, 1.0])
        ax1.set_ylim([0.0, 1.05])
        ax1.set_xlabel('False Positive Rate')
        ax1.set_ylabel('True Positive Rate')
        ax1.set_title('ROC Curve')
        ax1.legend(loc="lower right")

    y_pred = model.predict(X_test)
    cm = confusion_matrix(y_test, y_pred)
    disp = ConfusionMatrixDisplay(confusion_matrix=cm)
    disp.plot(ax=ax2, cmap=plt.cm.Reds)
    ax2.set_title('Confusion Matrix')

    plt.tight_layout()
    plt.show()

    acc = accuracy_score(y_test, y_pred)
    prec = precision_score(y_test, y_pred)
    rec = recall_score(y_test, y_pred)
    f1 = f1_score(y_test, y_pred)

    print(f"\n{name}")
    print("Accuracy:", acc)
    print("Precision:", prec)
    print("Recall:", rec)
    print("F1 Score:", f1)

```

```
    return acc, prec, rec, f1
```

5. Hyperparameter Spaces

```
# ================= HYPERPARAMETER SPACES =================
models_params = {
    "Decision Tree": (
        DecisionTreeClassifier(random_state=42),
        {"criterion": ["gini", "entropy"], "max_depth": [3, 5, 10, None]}
    ),

    "Random Forest": (
        RandomForestClassifier(random_state=42),
        {"n_estimators": [50, 100, 200], "max_depth": [3, 5, 10, None], "
    criterion": ["gini", "entropy"]}
    ),

    "AdaBoost": (
        AdaBoostClassifier(random_state=42, estimator=
    DecisionTreeClassifier(random_state=42)),
        {"n_estimators": [50, 100, 200], "learning_rate": [0.01, 0.1, 1],
         "estimator__max_depth": [1, 3, 5]}
    ),

    "Gradient Boosting": (
        GradientBoostingClassifier(random_state=42),
        {"n_estimators": [50, 100, 200], "learning_rate": [0.01, 0.1,
    0.5], "max_depth": [3, 5, 7]}
    ),

    "XGBoost": (
        xgb.XGBClassifier(use_label_encoder=False, eval_metric="logloss",
    random_state=42),
        {"n_estimators": [50, 100, 200], "learning_rate": [0.01, 0.1,
    0.3], "max_depth": [3, 5, 7],
         "gamma": [0, 0.1, 0.3]}
    ),
}

results_table = []
trial_tables = {}
best_estimators = {}
```

6. GridSearch and Top-5 Hyperparameter Trials

```
# ================= GRIDSEARCH + TRIAL TABLE =================
for name, (model, params) in models_params.items():
```

```
    print(f"\n--- Grid Search for {name} ---")
    grid = GridSearchCV(model, params, cv=5, scoring="accuracy", n_jobs=-1, return_train_score=False)
    grid.fit(X_train, y_train)

    best_model = grid.best_estimator_
    best_estimators[name] = best_model
    print("Best Params:", grid.best_params_)
    print("Best CV Score:", grid.best_score_)

    # Evaluate
    start = time.time()
    best_model.fit(X_train, y_train)
    elapsed = time.time() - start
    res = evaluate(name, best_model, X_test, y_test)
    results_table.append((name, grid.best_params_, *res, elapsed))

    # Collect top 5 hyperparameter trials
    trial_res = []
    for i in range(len(grid.cv_results_["params"])):
        trial_res.append({
            **grid.cv_results_["params"][i],
            "CV Accuracy": grid.cv_results_["mean_test_score"][i]
        })
    trial_df = pd.DataFrame(trial_res).sort_values(by="CV Accuracy", ascending=False).head(5)

    y_pred = best_model.predict(X_test)
    trial_df["F1 Score (Test)"] = f1_score(y_test, y_pred)

    trial_tables[name] = trial_df

    print(f"\nTop 5 Hyperparameter Trials for {name}")
    print(trial_df)
```

7. Stacking Classifiers

```
# ================= STACKING CLASSIFIER (3 Variants) =================
stacking_variants = {
    "Stacking (SVM+NB+DT     Logistic Regression)": StackingClassifier(
        estimators=[
            ("svm", SVC(probability=True, kernel="rbf", C=1, gamma="scale")),
            ("nb", GaussianNB()),
            ("dt", DecisionTreeClassifier(max_depth=5, random_state=42))
        ],
        final_estimator=LogisticRegression(max_iter=500, random_state=42)
    ),

    "Stacking (SVM+NB+DT     Random Forest)": StackingClassifier(
        estimators=[
```

```
            ("svm", SVC(probability=True, kernel="rbf", C=1, gamma="scale")),
            ("nb", GaussianNB()),
            ("dt", DecisionTreeClassifier(max_depth=5, random_state=42))
        ],
        final_estimator=RandomForestClassifier(n_estimators=100, random_state=42)
    ),

    "Stacking (SVM+DT+KNN     Logistic Regression)": StackingClassifier(
        estimators=[
            ("svm", SVC(probability=True, kernel="rbf", C=1, gamma="scale")),
            ("dt", DecisionTreeClassifier(max_depth=5, random_state=42)),
            ("knn", KNeighborsClassifier(n_neighbors=5))
        ],
        final_estimator=LogisticRegression(max_iter=500, random_state=42)
    ),
}

for name, stack_model in stacking_variants.items():
    start = time.time()
    stack_model.fit(X_train, y_train)
    elapsed = time.time() - start
    res = evaluate(name, stack_model, X_test, y_test)
    results_table.append((name, "Default (base learners tuned separately)", *res, elapsed))
    best_estimators[name] = stack_model
```

8. K-Fold Cross Validation

```
# ================ K-FOLD CROSS VALIDATION ================
print("\n--- 5-Fold Cross-Validation ---")
kf = KFold(n_splits=5, shuffle=True, random_state=42)
cv_results = {}

for name, model in best_estimators.items():
    scores = cross_val_score(model, X_scaled, y, cv=kf, scoring="accuracy")
    cv_results[name] = scores
    print(f"{name} Fold Accuracies: {scores}")
    print(f"{name} Avg Accuracy: {np.mean(scores):.4f}")
```

Results and Comparisons

Table 1: Model Performance with Tuned Hyperparameters

Model	Best Hyperparameters
Decision Tree	{criterion=entropy, max_depth=10}
Random Forest	{criterion=gini, max_depth=10, n_estimators=100}
AdaBoost	{max_depth=3, learning_rate=1, n_estimators=50}
Gradient Boosting	{learning_rate=0.5, max_depth=3, n_estimators=200}
XGBoost	{gamma=0, learning_rate=0.3, max_depth=3, n_estimators=200}
Stacking (SVM+NB+DT → LR)	Default (tuned base learners)
Stacking (SVM+NB+DT → RF)	Default (tuned base learners)
Stacking (SVM+DT+KNN → LR)	Default (tuned base learners)

Table 2: K-Fold CV Accuracies (K=5)

Fold	Decision Tree	Random Forest	AdaBoost	GB	XGBoost	Stacking (SVM+
1	0.9474	0.9561	0.9649	0.9649	0.9561	0.96
2	0.9561	0.9649	0.9737	1.0000	0.9649	0.99
3	0.9123	0.9386	0.9561	0.9474	0.9561	0.96
4	0.9474	0.9561	0.9912	0.9912	0.9737	0.98
5	0.9558	0.9646	0.9469	0.9381	0.9646	0.98
Average	0.9438	0.9561	0.9666	0.9683	0.9631	0.97

Table 3: Top-5 Hyperparameter Trials

Decision Tree

Criterion	Max Depth	CV Accuracy	F1 Score (Test)
entropy	10	0.9363	0.9383
entropy	None	0.9363	0.9383
gini	5	0.9341	0.9383
entropy	5	0.9341	0.9383
gini	3	0.9319	0.9383

Random Forest

Criterion	Max Depth	Estimators	CV Accuracy	F1 Score (Test)
gini	10	50	0.9670	0.9630
gini	None	50	0.9670	0.9630
entropy	None	50	0.9626	0.9630
gini	10	100	0.9626	0.9630
entropy	10	50	0.9626	0.9630

AdaBoost

Estimator Depth	Learning Rate	Estimators	CV Accuracy	F1 Score (Test)
3	1.0	50	0.9692	0.9500
3	0.1	100	0.9670	0.9500
1	1.0	200	0.9648	0.9500
3	1.0	200	0.9626	0.9500
3	1.0	100	0.9626	0.9500

Gradient Boosting

Learning Rate	Max Depth	Estimators	CV Accuracy	F1 Score (Test)
0.5	3	200	0.9582	0.9500
0.5	3	100	0.9582	0.9500
0.1	3	50	0.9560	0.9500
0.5	3	50	0.9538	0.9500
0.1	3	200	0.9538	0.9500

XGBoost

Gamma	Learning Rate	Max Depth	Estimators	CV Accuracy	F1 Score (Test)
0.0	0.3	3	200	0.9714	0.9630
0.0	0.3	3	50	0.9692	0.9630
0.0	0.3	3	100	0.9692	0.9630
0.0	0.3	7	100	0.9670	0.9630
0.0	0.3	7	200	0.9670	0.9630

ROC Curves and Confusion Matrices

Below are the ROC curves and confusion matrices for each model:

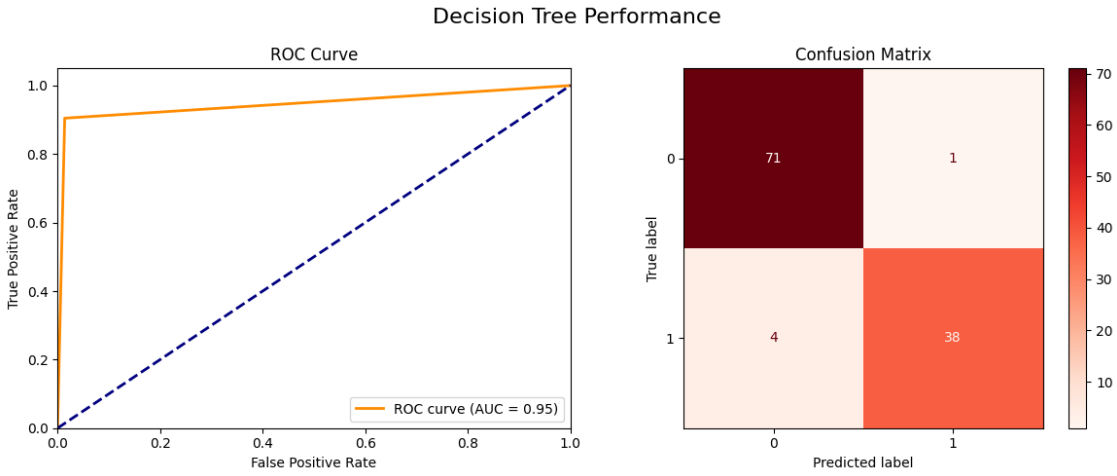

Figure 1: Decision Tree ROC + Confusion Matrix

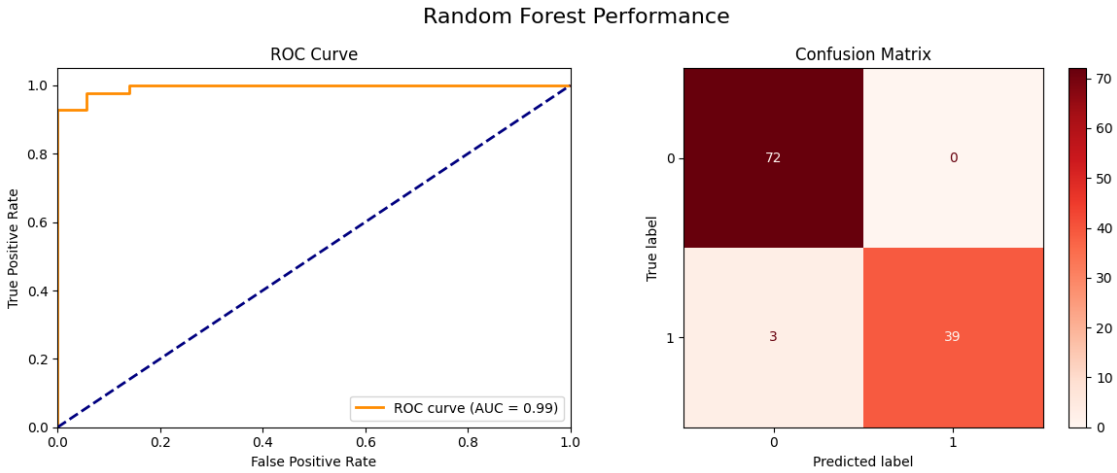

Figure 2: Random Forest ROC + Confusion Matrix

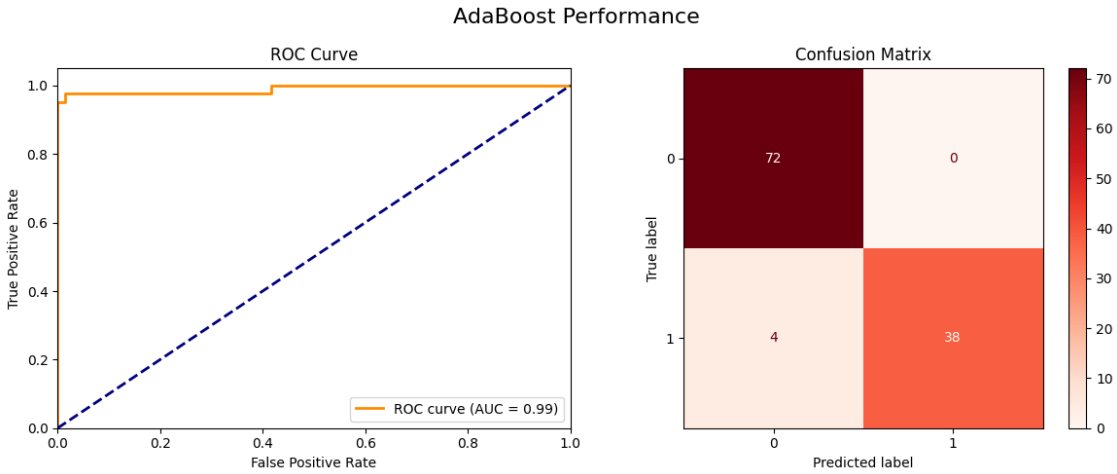

Figure 3: AdaBoost ROC + Confusion Matrix

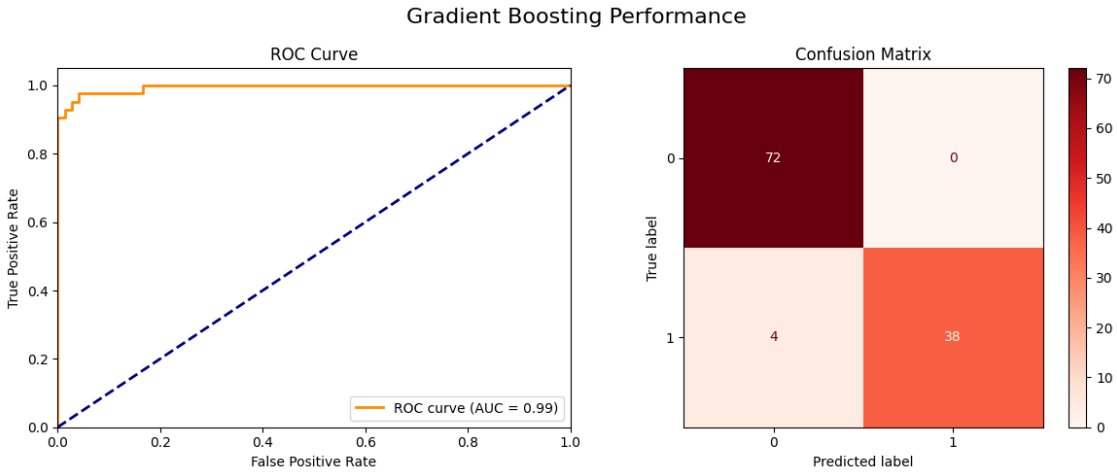

Figure 4: Gradient Boosting ROC + Confusion Matrix

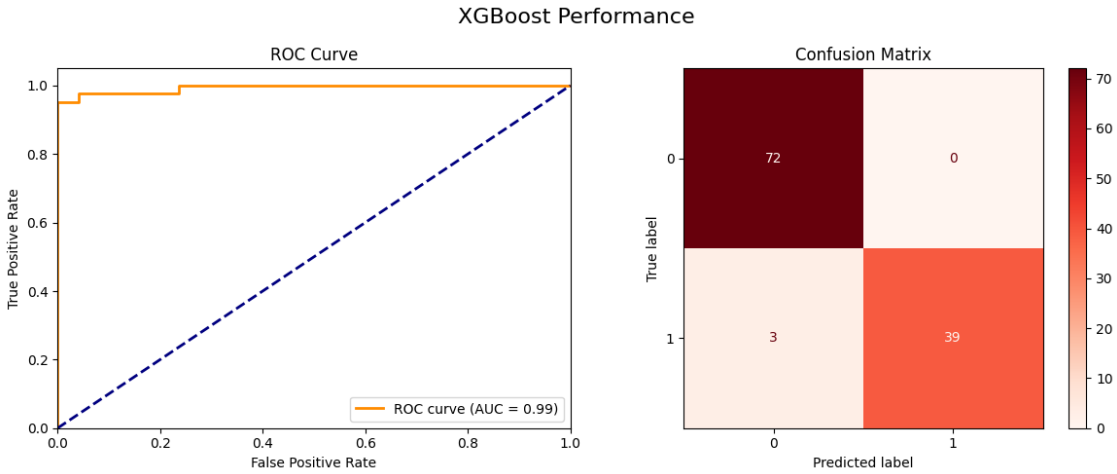

Figure 5: XGBoost ROC + Confusion Matrix

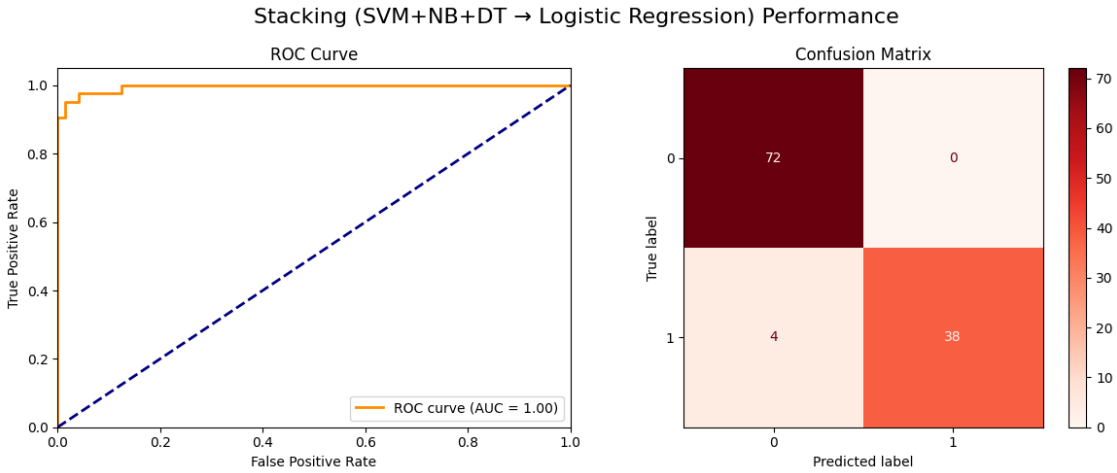

Figure 6: Stacking (SVM+NB+DT → Logistic Regression) ROC + Confusion Matrix

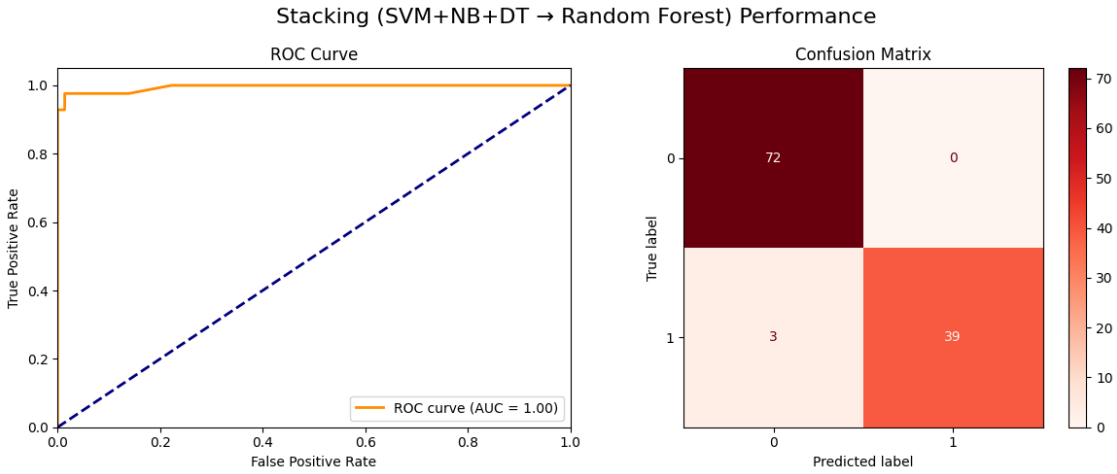

Figure 7: Stacking (SVM+NB+DT → Random Forest) ROC + Confusion Matrix

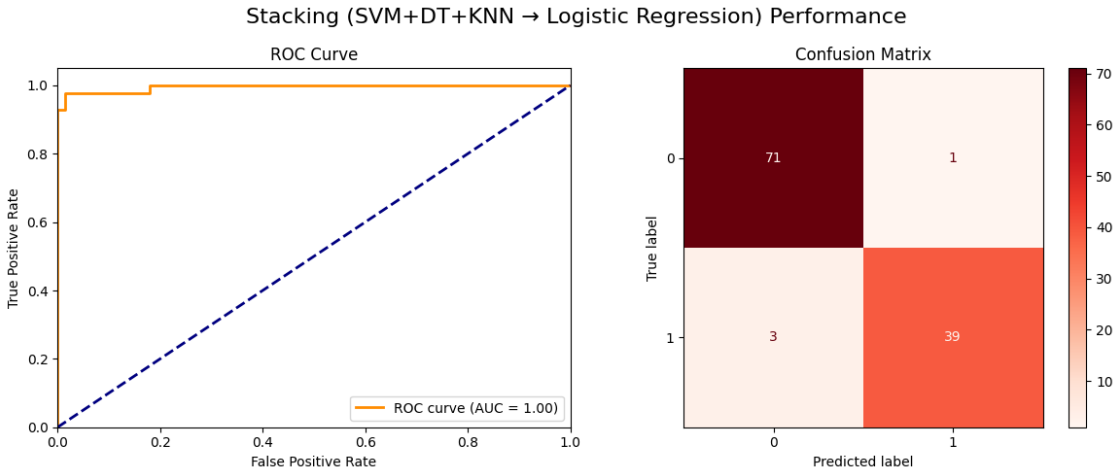

Figure 8: Stacking (SVM+DT+KNN → Logistic Regression) ROC + Confusion Matrix

Conclusion

In this experiment, multiple classifiers were optimized and evaluated on the Breast Cancer Wisconsin dataset:

- **Decision Tree:** Strong baseline, 94% CV accuracy, but slightly weaker recall.
- **Random Forest:** Improved generalization with 95.6% CV accuracy.
- **AdaBoost & Gradient Boosting:** Both performed well with ≈96–97% CV accuracy, showing boosting's strength.
- **XGBoost:** Achieved 96.3% CV accuracy and high F1 score, confirming its robustness.
- **Stacking:** Best results overall, especially Stacking (SVM+NB+DT → Logistic Regression) with 97.7% CV accuracy.

Thus, ensemble and stacking methods outperform single classifiers, with stacking providing the most reliable breast cancer classification results.