

Experiment 5: Perceptron Learning Algorithm (PLA)

One-vs-Rest Multi-class Classification

Prepared by: **Sreeram GM**
Course / Lab: Experiment 5

September 9, 2025

Contents

1	Aim	1
2	Dataset	1
3	Preprocessing	2
4	Implementation	2
4.1	Mounting and CSV / Image loading	2
4.2	Preprocessing function and dataset build	3
4.3	Label encoding and splitting	4
4.4	PLA (One-vs-Rest) implementation	4
4.5	Hyperparameter search (grid) and retraining	5
4.6	Evaluation and saving	6
5	Hyperparameter Search and Best Configuration	6
6	Results	6
6.1	Overall summary	6
6.2	Per-class precision / recall / f1	6
6.3	Figures (placeholders)	8
7	Discussion	9
8	Conclusion	10

1 Aim

Implement the Perceptron Learning Algorithm (PLA) as a one-vs-rest multi-class classifier for the image dataset, perform hyperparameter search, evaluate the best model on a held-out test set, and present results (metrics, confusion matrix, ROC, training-error curve).

2 Dataset

- Dataset root (Colab / Google Drive): `/content/drive/MyDrive/colabdata/dataset`
- Contents:
 - `Img/` : folder containing image files.

- *.csv : CSV file mapping image filenames to labels.
- Summary (from experiment run):
 - Number of classes: **62**
 - Train samples: **2728**
 - Test samples: **682**

3 Preprocessing

1. Mount Google Drive in Colab and locate dataset folder.
2. Load CSV mapping filenames to labels; auto-detect filename and label columns when possible.
3. For each image:
 - Convert to grayscale (PIL 'L').
 - Resize to **28×28**.
 - Normalize pixel values to [0, 1].
 - Flatten to a vector of length 784.
4. Use `LabelEncoder` to convert labels into integer class indices.
5. Use stratified train/test split (80% train, 20% test). During hyperparameter search, use an internal validation split from train.

4 Implementation

Below are the key code sections used in the experiment. The full code is included in the appendix.

4.1 Mounting and CSV / Image loading

```

1 from google.colab import drive
2 drive.mount('/content/drive')
3
4 DATASET_ROOT = '/content/drive/MyDrive/colabdata/dataset'
5 IMG_FOLDER = DATASET_ROOT + '/Img'
6
7 # find CSV file in dataset root
8 def find_csv_file(dataset_root):
9     candidates = [f for f in os.listdir(dataset_root) if f.lower().
10                    endswith('.csv')]
11     if not candidates:
12         raise FileNotFoundError(f"No CSV file found in {dataset_root}.
13                                Place CSV in that folder.")
14     return os.path.join(dataset_root, candidates[0])
15
16 CSV_PATH = find_csv_file(DATASET_ROOT)
17 df = pd.read_csv(CSV_PATH).dropna().reset_index(drop=True)
18
19 # auto-detect filename & label columns

```

```

18 possible_file_cols = ['filename', 'file', 'image', 'img', 'path', '
    image_path', 'file_name', 'File']
19 possible_label_cols = ['label', 'class', 'target', 'y', 'label_name', 'Label
    ']
20 file_col = None
21 label_col = None
22 for c in df.columns:
23     low = c.lower()
24     if low in possible_file_cols or any(p in low for p in
        possible_file_cols):
25         file_col = c
26     if low in possible_label_cols or any(p in low for p in
        possible_label_cols):
27         label_col = c
28 if file_col is None: file_col = df.columns[0]
29 if label_col is None: label_col = df.columns[1] if len(df.columns)>1
    else df.columns[0]
30 print("Detected file column:", file_col)
31 print("Detected label column:", label_col)

```

Listing 1: Mount Drive and auto-detect CSV + load filenames

4.2 Preprocessing function and dataset build

```

1 from PIL import Image
2 import numpy as np
3
4 def load_and_preprocess(img_path, size=(28,28), as_gray=True):
5     img = Image.open(img_path)
6     if as_gray:
7         img = img.convert('L')
8     img = img.resize(size, Image.BILINEAR)
9     arr = np.asarray(img, dtype=np.float32)/255.0
10    return arr.flatten()
11
12 # Resolve paths and build lists
13 image_paths = []
14 labels = []
15 missing_files = []
16 for idx, row in df.iterrows():
17     fname = str(row[file_col])
18     full_path = fname
19     if not os.path.isabs(full_path):
20         p1 = os.path.join(IMG_FOLDER, fname)
21         p2 = os.path.join(IMG_FOLDER, os.path.basename(fname))
22         if os.path.exists(p1): full_path = p1
23         elif os.path.exists(p2): full_path = p2
24     else:
25         p3 = os.path.join(DATASET_ROOT, fname)
26         if os.path.exists(p3): full_path = p3
27     else:
28         missing_files.append(fname)
29         continue
30     image_paths.append(full_path)
31     labels.append(row[label_col])
32
33 # load images (can be slow)

```

```

34 X_list = []
35 for p in tqdm(image_paths):
36     X_list.append(load_and_preprocess(p, size=(28,28)))
37 X = np.vstack(X_list)
38 y = np.array(labels)

```

Listing 2: Image preprocessing and building X,y arrays

4.3 Label encoding and splitting

```

1 from sklearn.preprocessing import LabelEncoder
2 from sklearn.model_selection import train_test_split
3
4 le = LabelEncoder()
5 y_enc = le.fit_transform(y)
6 classes = le.classes_
7 n_classes = len(classes)
8
9 X_train, X_test, y_train_idx, y_test_idx = train_test_split(
10     X, y_enc, test_size=0.2, stratify=y_enc, random_state=42)

```

Listing 3: Label encoding and train/test split

4.4 PLA (One-vs-Rest) implementation

```

1 class PerceptronPLA:
2     def __init__(self, n_features, lr=1.0):
3         self.lr = lr
4         self.w = np.zeros(n_features + 1, dtype=np.float32) # include
                    bias
5
6     def predict_raw(self, X):
7         Xb = np.hstack([X, np.ones((X.shape[0],1), dtype=np.float32)])
8         return Xb.dot(self.w)
9
10    def predict(self, X):
11        return np.where(self.predict_raw(X) >= 0, 1, -1)
12
13    def fit(self, X, y, epochs=10, shuffle=True, verbose=False):
14        Xb = np.hstack([X, np.ones((X.shape[0],1), dtype=np.float32)])
15        n = X.shape[0]
16        history = []
17        for ep in range(epochs):
18            errors = 0
19            indices = np.arange(n)
20            if shuffle: np.random.shuffle(indices)
21            for i in indices:
22                xi = Xb[i]; yi = y[i]
23                pred = 1 if xi.dot(self.w) >= 0 else -1
24                if pred != yi:
25                    self.w += self.lr * yi * xi
26                    errors += 1
27            history.append(errors / n)
28        return history
29
30    def train_ovr_pla(X_train, y_train, lr=1.0, epochs=20):

```

```

31     n_features = X_train.shape[1]
32     models = {}
33     history = {}
34     for c_idx, c_label in enumerate(classes):
35         y_bin = np.where(y_train == c_idx, 1, -1)
36         p = PerceptronPLA(n_features, lr=lr)
37         hist = p.fit(X_train, y_bin, epochs=epochs)
38         models[c_idx] = p
39         history[c_idx] = hist
40     return models, history

```

Listing 4: Perceptron PLA class and OvR training

4.5 Hyperparameter search (grid) and retraining

```

1  # grid
2  lr_values = [1.0, 0.1, 0.01]
3  epoch_values = [10, 20, 50]
4
5  best_acc = 0.0
6  best_params = None
7  best_models = None
8  best_history = None
9
10 search_X_tr, search_X_val, search_y_tr, search_y_val = train_test_split
    (
11     X_train, y_train_idx, test_size=0.2, stratify=y_train_idx,
        random_state=42)
12
13 for lr in lr_values:
14     for ep in epoch_values:
15         candidate_models, candidate_hist = train_ovr_pla(search_X_tr,
16                 search_y_tr, lr=lr, epochs=ep)
17         # validation predictions
18         def ovr_predict_local(models, X_input):
19             n_samples = X_input.shape[0]
20             scores_loc = np.zeros((n_samples, len(models)), dtype=np.
21                 float32)
22             for c_idx, m in models.items():
23                 scores_loc[:, c_idx] = m.predict_raw(X_input)
24             preds_loc = np.argmax(scores_loc, axis=1)
25             return preds_loc
26         val_preds = ovr_predict_local(candidate_models, search_X_val)
27         acc = accuracy_score(search_y_val, val_preds)
28         if acc > best_acc:
29             best_acc = acc
30             best_params = (lr, ep)
31             best_models = candidate_models
32             best_history = candidate_hist
33
34 # retrain on full X_train with best params
35 best_lr, best_ep = best_params
36 models, hist = train_ovr_pla(X_train, y_train_idx, lr=best_lr, epochs=
    best_ep)

```

Listing 5: Grid search over LR and epochs; retrain best on full train

4.6 Evaluation and saving

```
1 def ovr_predict(models, X):
2     n = X.shape[0]
3     scores = np.zeros((n, len(models)), dtype=np.float32)
4     for c_idx, model in models.items():
5         scores[:, c_idx] = model.predict_raw(X)
6     preds = np.argmax(scores, axis=1)
7     return preds, scores
8
9 y_pred_idx, raw_scores_test = ovr_predict(models, X_test)
10
11 # metrics
12 acc = accuracy_score(y_test_idx, y_pred_idx)
13 prec, rec, f1, _ = precision_recall_fscore_support(y_test_idx,
14     y_pred_idx, average='macro', zero_division=0)
15
16 # save models
17 with open('/content/pla_ovr_models.pkl', 'wb') as f:
18     pickle.dump({'best_params': best_params, 'models': {c:m.w for c,m in
19         models.items()}}, 'label_encoder': le}, f)
```

Listing 6: Prediction, metrics, plots, save weights

5 Hyperparameter Search and Best Configuration

- Grid searched: Learning rates = {1.0, 0.1, 0.01}; Epochs = {10, 20, 50}.
- Best hyperparameters: **learning rate = 0.1, epochs = 50**.
- Best validation accuracy (internal validation): **0.1978**.

6 Results

6.1 Overall summary

- Number of classes: **62**
- Train samples: **2728**
- Test samples: **682**
- Best LR, epochs: **(0.1, 50)**
- Validation accuracy (best): **0.1978**
- Accuracy (test): **0.1833**
- Macro F1: **0.1674**

6.2 Per-class precision / recall / f1

Table 1: Per-class metrics (test set)

class_label	precision	recall	f1
0	0.214286	0.272727	0.240000
1	0.083969	1.000000	0.154930
2	1.000000	0.090909	0.166667
3	0.666667	0.181818	0.285714
4	0.000000	0.000000	0.000000
5	0.500000	0.090909	0.153846
6	0.000000	0.000000	0.000000
7	0.000000	0.000000	0.000000
8	0.666667	0.181818	0.285714
9	0.000000	0.000000	0.000000
A	0.000000	0.000000	0.000000
B	1.000000	0.272727	0.428571
C	0.250000	0.727273	0.372093
D	1.000000	0.090909	0.166667
E	0.375000	0.272727	0.315789
F	0.000000	0.000000	0.000000
G	0.444444	0.363636	0.400000
H	1.000000	0.090909	0.166667
I	0.000000	0.000000	0.000000
J	0.375000	0.545455	0.444444
K	0.000000	0.000000	0.000000
L	1.000000	0.363636	0.533333
M	0.666667	0.363636	0.470588
N	1.000000	0.090909	0.166667
O	0.166667	0.090909	0.117647
P	0.470588	0.727273	0.571429
Q	0.291667	0.636364	0.400000
R	1.000000	0.090909	0.166667
S	0.000000	0.000000	0.000000
T	0.454545	0.454545	0.454545
U	0.600000	0.272727	0.375000
V	0.363636	0.363636	0.363636
W	0.000000	0.000000	0.000000
X	1.000000	0.181818	0.307692
Y	0.000000	0.000000	0.000000
Z	0.500000	0.090909	0.153846
a	1.000000	0.090909	0.166667
b	0.000000	0.000000	0.000000
c	0.333333	0.181818	0.235294
d	0.133333	0.181818	0.153846
e	0.076923	0.090909	0.083333
f	0.000000	0.000000	0.000000
g	0.111111	0.272727	0.157895
h	0.045455	0.272727	0.077922
i	0.333333	0.181818	0.235294
j	0.200000	0.090909	0.125000
k	0.014925	0.090909	0.025641
l	0.000000	0.000000	0.000000
m	1.000000	0.090909	0.166667
n	0.000000	0.000000	0.000000
o	0.000000	0.000000	0.000000
p	0.571429	0.363636	0.444444
q	0.114286	0.727273	0.197531
r	0.153846	0.363636	0.216216

class_label	precision	recall	f1
s	0.130435	0.272727	0.176471
t	0.083333	0.090909	0.086957
u	0.000000	0.000000	0.000000
v	0.000000	0.000000	0.000000
w	0.000000	0.000000	0.000000
x	0.000000	0.000000	0.000000
y	1.000000	0.090909	0.166667
z	0.000000	0.000000	0.000000

6.3 Figures (placeholders)




Figure 1: Confusion matrix (replace with actual image path).

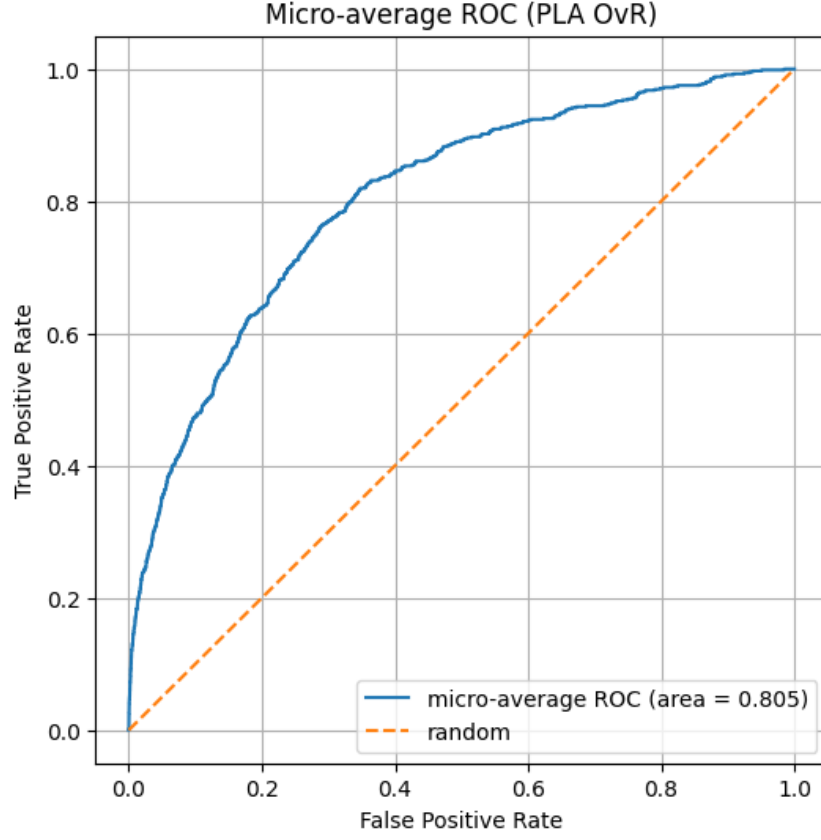


Figure 2: Micro-average ROC curve (replace with actual image path).



Figure 3: Training error vs epochs (average over classes).

7 Discussion

- Test accuracy is low (about 18.33%) and macro F1 is 0.1674. This indicates the dataset is challenging for PLA, likely due to:

- Large number of classes (62) with limited samples per class.
- High intra-class variability and possible class imbalance.
- PLA is a linear classifier and may not capture complex patterns in image data.
- Some classes have very high precision but very low recall (or vice versa), indicating the classifier is biased / sparse in predictions for those classes.
- Consider stronger models (MLP/CNN), data augmentation, feature extraction, or dimensionality reduction for improved performance.

8 Conclusion

The PLA (OvR) baseline was implemented and tuned with grid search. While it works end-to-end and provides interpretable per-class metrics, its performance on this image classification task is limited. Use the results as a baseline for comparison with MLP/CNN models.