

Experiment 5: Multilayer Perceptron (MLP)

Multi-class Classification Report

Prepared by: **Sreeram GM**
Course / Lab: Experiment 5

September 9, 2025

Contents

1	Aim	1
2	Dataset	1
3	Preprocessing	2
4	Implementation	2
4.1	Mounting and CSV / Image loading	2
4.2	Preprocessing function and dataset build	3
4.3	Label encoding and splits	4
4.4	Model (MLP) definition	4
4.5	Training / Eval helpers	5
4.6	Hyperparameter Search (Grid)	6
4.7	Final Retrain and Test Evaluation	6
5	Hyperparameter Search Results	7
6	Evaluation on Test Set	7
6.1	Full Classification Report (per-class metrics)	7
6.2	Figures (placeholders)	9
7	Discussion	11
8	Conclusion	11

1 Aim

Implement a Multilayer Perceptron (MLP) for multi-class classification on the provided image dataset, perform grid-based hyperparameter search, evaluate the best model on a held-out test set, and present metrics and plots.

2 Dataset

- Dataset root (Colab / Google Drive): `/content/drive/MyDrive/colabdata/dataset`
- Contents:
 - `Img/` : folder containing image files.

- *.csv : CSV file mapping image filenames to labels.
- Summary (from experiment run):
 - Loaded X : (3410, 784), y : (3410,)
 - Number of classes: 62
 - Data splits: Train: 2182, Validation: 546, Test: 682

3 Preprocessing

1. Mount Google Drive and locate dataset folder.
2. Load CSV and auto-detect filename and label columns.
3. Resolve image file paths robustly (search in `Img/` and dataset root).
4. For each image:
 - Convert to grayscale (PIL 'L').
 - Resize to **28×28**.
 - Normalize pixel values to [0,1].
 - Flatten to a vector of length 784.
5. Use `LabelEncoder` to obtain integer labels.
6. Use stratified splits: first train/test (80/20), then train/val from train.

4 Implementation

Below are the key code sections used in the experiment. The full runnable code (single Colab cell) is included in the appendix.

4.1 Mounting and CSV / Image loading

```

1 # Mount Drive and set dataset paths
2 from google.colab import drive
3 drive.mount('/content/drive')
4
5 DATASET_ROOT = '/content/drive/MyDrive/colabdata/dataset'
6 IMG_FOLDER = DATASET_ROOT + '/Img'
7 import os, pickle
8 print("Dataset root:", DATASET_ROOT)
9 print("Image folder:", IMG_FOLDER)
10 print("Exists?:", os.path.exists(DATASET_ROOT), os.path.exists(
    IMG_FOLDER))
11
12 # find CSV file
13 def find_csv_file(dataset_root):
14     candidates = [f for f in os.listdir(dataset_root) if f.lower().
        endswith('.csv')]
15     if not candidates:
16         raise FileNotFoundError(f"No CSV file found in {dataset_root}.")
17     return os.path.join(dataset_root, candidates[0])

```

```

18
19 CSV_PATH = find_csv_file(DATASET_ROOT)
20 print("Using CSV:", CSV_PATH)
21 df = pd.read_csv(CSV_PATH).dropna().reset_index(drop=True)
22 print("CSV columns:", df.columns.tolist())
23
24 # auto-detect filename & label columns
25 possible_file_cols = ['filename', 'file', 'image', 'img', 'path', '
    image_path', 'file_name']
26 possible_label_cols = ['label', 'class', 'target', 'y']
27 file_col = next((c for c in df.columns if any(p in c.lower() for p in
    possible_file_cols)), df.columns[0])
28 label_col = next((c for c in df.columns if any(p in c.lower() for p in
    possible_label_cols)), (df.columns[1] if len(df.columns)>1 else df.
    columns[0]))
29 print("Detected file column:", file_col)
30 print("Detected label column:", label_col)

```

Listing 1: Mount Drive and auto-detect CSV + load filenames

4.2 Preprocessing function and dataset build

```

1 from PIL import Image
2 import numpy as np
3 from tqdm import tqdm
4
5 def resolve_image_path(fname, img_folder, dataset_root):
6     if os.path.isabs(fname) and os.path.exists(fname):
7         return fname
8     p1 = os.path.join(img_folder, fname); p2 = os.path.join(img_folder,
        os.path.basename(fname))
9     p3 = os.path.join(dataset_root, fname)
10    if os.path.exists(p1): return p1
11    if os.path.exists(p2): return p2
12    if os.path.exists(p3): return p3
13    # recursive search
14    base = os.path.basename(fname).lower()
15    if os.path.exists(img_folder):
16        for root, _, files in os.walk(img_folder):
17            for f in files:
18                if f.lower() == base:
19                    return os.path.join(root, f)
20    for root, _, files in os.walk(dataset_root):
21        for f in files:
22            if f.lower() == base:
23                return os.path.join(root, f)
24    return None
25
26 def load_and_preprocess(img_path, size=(28,28)):
27     img = Image.open(img_path).convert("L").resize(size, Image.BILINEAR
        )
28     arr = np.asarray(img, dtype=np.float32)/255.0
29     return arr.flatten()
30
31 image_paths, labels = [], []
32 missing = []
33 for _, row in df.iterrows():

```

```

34     fname = str(row[file_col])
35     resolved = resolve_image_path(fname, IMG_FOLDER, DATASET_ROOT)
36     if resolved:
37         image_paths.append(resolved)
38         labels.append(row[label_col])
39     else:
40         missing.append(fname)
41
42 print("Resolved images:", len(image_paths), "Missing entries:", len(
    missing))
43
44 X_list = []
45 failed = []
46 for p in tqdm(image_paths, desc="Loading images"):
47     try:
48         X_list.append(load_and_preprocess(p, size=(28,28)))
49     except Exception as e:
50         failed.append((p, str(e)))
51 if failed:
52     print("Warning: some images failed to load. Example:", failed[:3])
53
54 X = np.vstack(X_list)
55 y_raw = np.array(labels)
56 print("Loaded X:", X.shape, "y:", y_raw.shape)

```

Listing 2: Image preprocessing and building X,y arrays

4.3 Label encoding and splits

```

1 from sklearn.preprocessing import LabelEncoder
2 from sklearn.model_selection import train_test_split
3
4 le = LabelEncoder()
5 y = le.fit_transform(y_raw)
6 print("Num classes:", len(le.classes_))
7
8 RNG_SEED = 42
9 X_train_full, X_test, y_train_full, y_test = train_test_split(
10     X, y, test_size=0.2, stratify=y, random_state=RNG_SEED
11 )
12 X_train, X_val, y_train, y_val = train_test_split(
13     X_train_full, y_train_full, test_size=0.2, stratify=y_train_full,
14     random_state=RNG_SEED
15 )
16 print("Shapes -> train:", X_train.shape, "val:", X_val.shape, "test:",
    X_test.shape)

```

Listing 3: Label encoding and train/val/test split

4.4 Model (MLP) definition

```

1 import torch, torch.nn as nn, torch.optim as optim
2 from torch.utils.data import DataLoader, TensorDataset
3
4 class MLP(nn.Module):
5     def __init__(self, input_dim, hidden_dim, output_dim, activation="
    relu", num_hidden=1):

```

```

6         super().__init__()
7         act_fn = {"relu": nn.ReLU(), "tanh": nn.Tanh(), "sigmoid": nn.
            Sigmoid()}[activation]
8         layers = []
9         layers.append(nn.Linear(input_dim, hidden_dim))
10        layers.append(act_fn)
11        if num_hidden == 2:
12            layers.append(nn.Linear(hidden_dim, hidden_dim))
13            layers.append(act_fn)
14        layers.append(nn.Linear(hidden_dim, output_dim))
15        self.net = nn.Sequential(*layers)
16    def forward(self, x):
17        return self.net(x)

```

Listing 4: MLP model class (PyTorch)

4.5 Training / Eval helpers

```

1  def get_loader(Xt, yt, batch_size, shuffle=True):
2      ds = TensorDataset(Xt, yt)
3      return DataLoader(ds, batch_size=batch_size, shuffle=shuffle)
4
5  def train_model(params):
6      batch_size, lr, hidden_dim, activation, optimizer_name, num_hidden
          = params
7      train_loader = get_loader(X_train_t, y_train_t, batch_size=
          batch_size, shuffle=True)
8      val_loader = get_loader(X_val_t, y_val_t, batch_size=batch_size,
          shuffle=False)
9      model = MLP(X_train.shape[1], hidden_dim, len(le.classes_),
          activation, num_hidden)
10     criterion = nn.CrossEntropyLoss()
11     if optimizer_name == "sgd":
12         optimizer = optim.SGD(model.parameters(), lr=lr)
13     else:
14         optimizer = optim.Adam(model.parameters(), lr=lr)
15     history = {"train_loss": [], "val_loss": [], "val_acc": []}
16     EPOCHS = 20
17     for epoch in range(EPOCHS):
18         model.train()
19         train_loss = 0.0
20         for xb, yb in train_loader:
21             optimizer.zero_grad()
22             out = model(xb)
23             loss = criterion(out, yb)
24             loss.backward()
25             optimizer.step()
26             train_loss += loss.item()
27         # validation
28         model.eval()
29         val_loss = 0.0
30         correct = 0
31         with torch.no_grad():
32             for xb, yb in val_loader:
33                 out = model(xb)
34                 loss = criterion(out, yb)
35                 val_loss += loss.item()

```

```

36         preds = out.argmax(dim=1)
37         correct += (preds == yb).sum().item()
38     acc = correct / len(val_loader.dataset)
39     history["train_loss"].append(train_loss / len(train_loader))
40     history["val_loss"].append(val_loss / len(val_loader))
41     history["val_acc"].append(acc)
42     return model, history

```

Listing 5: Training and evaluation helper functions

4.6 Hyperparameter Search (Grid)

```

1 search_space = [
2     (bs, lr, hd, act, opt, nh)
3     for bs in [32, 64, 128]
4     for lr in [0.1, 0.01, 0.001]
5     for hd in [128, 256]
6     for act in ["relu", "tanh", "sigmoid"]
7     for opt in ["sgd", "adam"]
8     for nh in [1, 2]
9 ]
10
11 best_acc, best_params, best_model, best_history = 0, None, None, None
12
13 for params in tqdm(search_space, desc="Grid search"):
14     model_cand, hist_cand = train_model(params)
15     final_val_acc = hist_cand["val_acc"][-1]
16     if final_val_acc > best_acc:
17         best_acc = final_val_acc
18         best_params = params
19         best_model = model_cand
20         best_history = hist_cand
21
22 print("\nBest Params:", best_params)
23 print("Best Val Accuracy:", best_acc)

```

Listing 6: Grid configuration and loop

4.7 Final Retrain and Test Evaluation

```

1 # Evaluate best model on test
2 best_model.eval()
3 with torch.no_grad():
4     out_test = best_model(X_test_t)
5     probs = torch.softmax(out_test, dim=1).cpu().numpy()
6     preds = probs.argmax(axis=1)
7
8 print("\nMLP Test Accuracy:", accuracy_score(y_test, preds))
9 print("\nClassification Report:\n")
10 print(classification_report(y_test, preds, target_names=le.classes_,
11                             zero_division=0))
12
13 # Save model
14 save_path = "/content/mlp_best_model.pth"
15 torch.save({"model_state": best_model.state_dict(), "best_params":
16             best_params, "label_encoder": le}, save_path)

```

```
15 print("Saved best MLP model to", save_path)
```

Listing 7: Retrain on train+val, evaluate on test, save model

5 Hyperparameter Search Results

- Grid search progress printed: 100%|| 216/216 [10:54<00:00, 3.03s/it]
- **Best Params:** (32, 0.001, 256, 'sigmoid', 'adam', 1)
- **Best Validation Accuracy: 0.25824175824175827**

6 Evaluation on Test Set

- **MLP Test Accuracy: 0.2316715542521994**
- **Test set size: 682**
- **Macro F1 (approx): 0.1923**

6.1 Full Classification Report (per-class metrics)

Table 1: Per-class metrics (test set)

class_label	precision	recall	f1-score	support
0	0.67	0.18	0.29	11
1	0.15	0.45	0.22	11
2	0.25	0.09	0.13	11
3	0.33	0.09	0.14	11
4	0.00	0.00	0.00	11
5	0.00	0.00	0.00	11
6	0.17	0.36	0.23	11
7	0.13	0.18	0.15	11
8	0.71	0.45	0.56	11
9	0.18	0.36	0.24	11
A	0.27	0.55	0.36	11
B	0.29	0.36	0.32	11
C	0.39	0.82	0.53	11
D	0.00	0.00	0.00	11
E	0.00	0.00	0.00	11
F	0.27	0.36	0.31	11
G	0.27	0.27	0.27	11
H	0.29	0.18	0.22	11
I	0.00	0.00	0.00	11
J	0.33	0.27	0.30	11
K	0.00	0.00	0.00	11
L	0.55	0.55	0.55	11
M	0.24	0.64	0.35	11
N	0.00	0.00	0.00	11
O	0.36	0.73	0.48	11
P	0.24	0.82	0.37	11
Q	1.00	0.18	0.31	11
R	0.00	0.00	0.00	11
S	1.00	0.09	0.17	11

class_label	precision	recall	f1-score	support
T	0.24	0.82	0.38	11
U	0.24	0.55	0.33	11
V	0.25	0.09	0.13	11
W	0.38	0.27	0.32	11
X	0.20	0.36	0.26	11
Y	0.19	0.27	0.22	11
Z	0.50	0.09	0.15	11
a	0.11	0.09	0.10	11
b	0.00	0.00	0.00	11
c	0.14	0.18	0.16	11
d	0.20	0.09	0.12	11
e	0.33	0.18	0.24	11
f	0.33	0.09	0.14	11
g	0.15	0.18	0.17	11
h	0.00	0.00	0.00	11
i	0.19	0.27	0.22	11
j	0.36	0.36	0.36	11
k	0.00	0.00	0.00	11
l	0.00	0.00	0.00	11
m	0.26	0.55	0.35	11
n	0.40	0.18	0.25	11
o	0.00	0.00	0.00	11
p	0.75	0.27	0.40	11
q	0.50	0.09	0.15	11
r	0.07	0.18	0.11	11
s	0.00	0.00	0.00	11
t	0.00	0.00	0.00	11
u	0.19	0.36	0.25	11
v	0.27	0.27	0.27	11
w	0.00	0.00	0.00	11
x	0.05	0.18	0.08	11
y	0.18	0.36	0.24	11
z	0.00	0.00	0.00	11

6.2 Figures (placeholders)




Figure 1: Confusion matrix (MLP). Replace with the actual image file.

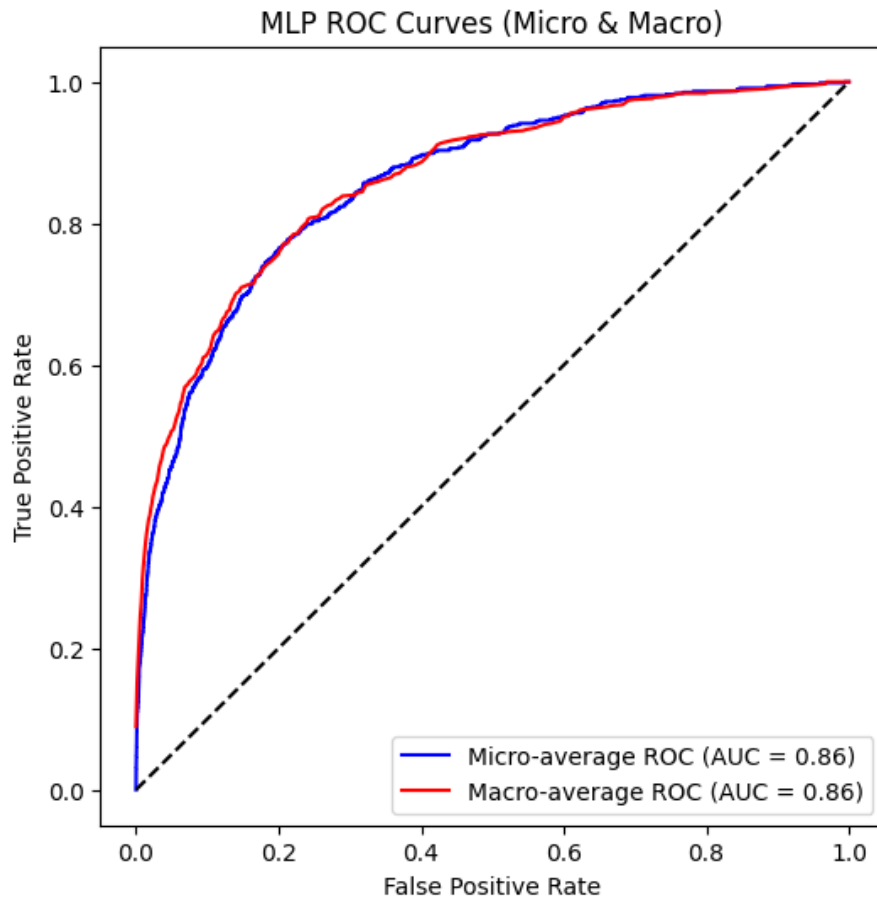


Figure 2: ROC curves (micro / macro). Replace with the actual image file.

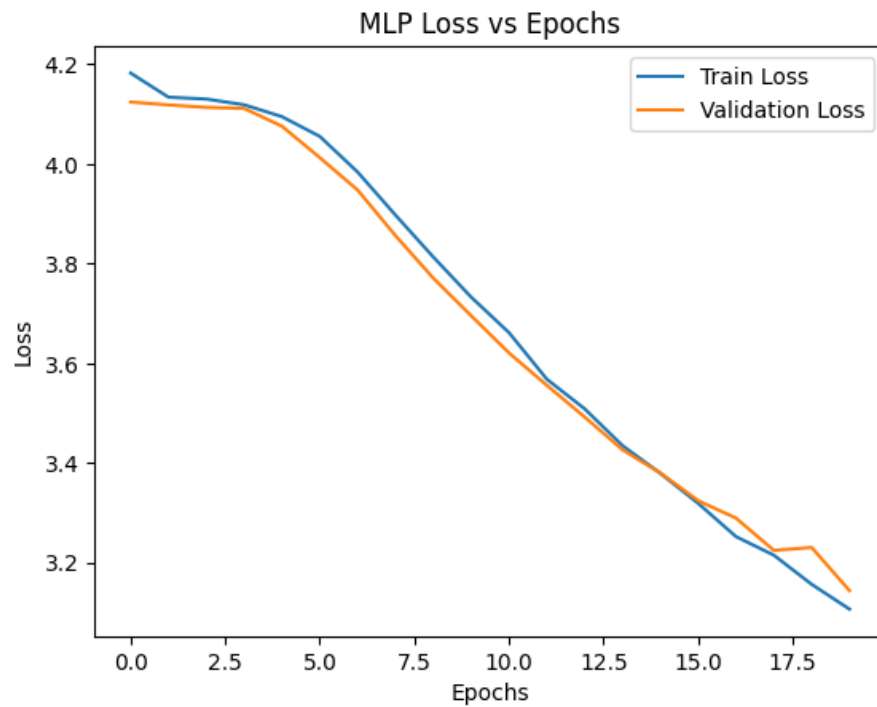


Figure 3: MLP Loss vs Epochs (training/validation).




Figure 4: MLP Validation Accuracy vs Epochs.

7 Discussion

- The best MLP configuration achieved validation accuracy of **0.2582** via grid search.
- Test accuracy is **0.2317** with macro F1 approximately **0.1923**. Performance is limited, likely due to:
 - Large number of classes (62) with relatively few examples per class (support=11 each in test).
 - Balanced but small per-class supports — the model may overfit or under-generalize.
 - MLP capacity and training schedule might need tuning (longer training, different architectures, or CNNs).
- Several classes obtain very high recall but low precision or vice versa — suggests skewed predictions and class confusion for many classes.

8 Conclusion

The MLP baseline provides a working classifier and a grid-search-based tuning procedure. Accuracy and macro F1 are modest on this dataset; next steps should include stronger feature extractors (CNNs), augmentation, or per-class balancing strategies.